

Python R training course - Pandas

Van-Duyet LE, Thinh PHAM

August 17, 2017

1 Introduction

Pandas is an open-source Python Library providing high-performance data manipulation and analysis tool using its powerful data structures. The name Pandas is derived from the word Panel Data – an Econometrics from Multidimensional data.

Python with Pandas is used in a wide range of fields including academic and commercial domains including finance, economics, Statistics, analytics, etc.

Key Features of Pandas

- Fast and efficient DataFrame object with default and customized indexing.
- Tools for loading data into in-memory data objects from different file formats.
- Data alignment and integrated handling of missing data.
- Reshaping and pivoting of date sets.
- Label-based slicing, indexing and subsetting of large data sets.
- Columns from a data structure can be deleted or inserted.
- Group by data for aggregation and transformations.
- High performance merging and joining of data.
- Time Series functionality.

2 Introduction to Data Structures

Pandas deals with the following three data structures:

- DataFrame
- Series
- Panel

These data structures are built on top of Numpy array, which means they are fast.

3 Python Pandas - DataFrame

A Data frame is a two-dimensional data structure, i.e., data is aligned in a tabular fashion in rows and columns.

3.1 Create DataFrame

3.1.1 Create an Empty DataFrame

```
import pandas as pd
df = pd.DataFrame()
print df
```

3.1.2 Create a DataFrame from Lists

```
In [4]: data = [1,2,3,4,5]
        df = pd.DataFrame(data)
        df
```

```
Out[4]:    0
0      1
1      2
2      3
3      4
4      5
```

```
In [9]: data = [['Duyet',10], ['Thinh',12], ['Nam',13]]
        df = pd.DataFrame(data,columns=['Name', 'Age'])
        df
```

```
Out[9]:    Name  Age
0  Duyet   10
1  Thinh   12
2    Nam   13
```

```
In [10]: data = [['Duyet',10], ['Thinh',12], ['Nam',13]]
        df = pd.DataFrame(data,columns=['Name', 'Age'], dtype=float)
        df
```

```
Out[10]:    Name  Age
0  Duyet  10.0
1  Thinh  12.0
2    Nam  13.0
```

3.1.3 Create a DataFrame from Dict of ndarrays / Lists

All the ndarrays must be of same length. If index is passed, then the length of the index should equal to the length of the arrays.

If no index is passed, then by default, index will be range(n), where n is the array length.

```
In [11]: data = {'Name':['Tom', 'Jack', 'Steve', 'Ricky'],'Age':[28,34,29,42]}
        df = pd.DataFrame(data)
        df
```

```
Out[11]:
```

	Age	Name
0	28	Tom
1	34	Jack
2	29	Steve
3	42	Ricky

3.1.4 Create a DataFrame from List of Dicts

```
In [14]: data = [{ 'a': 1, 'b': 2},{ 'a': 5, 'b': 10, 'c': 20}]
df = pd.DataFrame(data)
df

# PS: NaN (Not a Number) is appended in missing areas.
```

```
Out[14]:
```

	a	b	c
0	1	2	NaN
1	5	10	20.0

3.1.5 Reading from file (CSV, Excel, HDFS, SQL, ...)

Using `pd.read_csv()` function to read dataframe from a CSV file.

```
pandas.read_csv(filepath_or_buffer, sep='\t', delimiter=None, header='infer',
names=None, index_col=None, usecols=None)
```

Using `pd.read_excel()` to read Excel file.

```
pd.read_excel('foo.xlsx', 'Sheet1', index_col=None, na_values=['NA'])
```

Example: Here is how the csv file data (*data.csv*) looks like:

```
S.No,Name,Age,City,Salary
1,Tom,28,Toronto,20000
2,Lee,32,HongKong,3000
3,Steven,43,Bay Area,8300
4,Ram,38,Hyderabad,3900
```

```
In [6]: df=pd.read_csv("data.csv")
df
```

```
Out[6]:
```

	S.No	Name	Age	City	Salary
0	1	Tom	28	Toronto	20000
1	2	Lee	32	HongKong	3000
2	3	Steven	43	Bay Area	8300
3	4	Ram	38	Hyderabad	3900

Specify the names of the header using the **names** argument.

```
In [17]: df=pd.read_csv("data.csv", names=['a', 'b', 'c','d','e'])
df
```

```
Out [17]:
```

	a	b	c	d	e
0	S.No	Name	Age	City	Salary
1	1	Tom	28	Toronto	20000
2	2	Lee	32	HongKong	3000
3	3	Steven	43	Bay Area	8300
4	4	Ram	38	Hyderabad	3900

3.2 Viewing Data

```
In [89]: df.head()
```

```
Out [89]:
```

	S.No	Name	Age	City	Salary	Salary2
0	1	Tom	28	Toronto	20000	141.421356
1	2	Lee	32	HongKong	3000	54.772256
2	3	Steven	43	Bay Area	8300	91.104336
3	4	Ram	38	Hyderabad	3900	62.449980

```
In [91]: df.tail(3)
```

```
Out [91]:
```

	S.No	Name	Age	City	Salary	Salary2
1	2	Lee	32	HongKong	3000	54.772256
2	3	Steven	43	Bay Area	8300	91.104336
3	4	Ram	38	Hyderabad	3900	62.449980

```
In [92]: df.T
```

```
Out [92]:
```

	0	1	2	3
S.No	1	2	3	4
Name	Tom	Lee	Steven	Ram
Age	28	32	43	38
City	Toronto	HongKong	Bay Area	Hyderabad
Salary	20000	3000	8300	3900
Salary2	141.421	54.7723	91.1043	62.45

3.3 Column Selection

We will understand this by selecting a column from the DataFrame.

```
In [18]: df=pd.read_csv("data.csv")
df["Name"]
```

```
Out [18]:
```

0	Tom
1	Lee
2	Steven
3	Ram

Name: Name, dtype: object

```
In [20]: # Or this
df.Name
```

```
Out[20]: 0      Tom
         1      Lee
         2    Steven
         3      Ram
         Name: Name, dtype: object
```

3.4 Sort

```
In [94]: df.sort_values(by='Age')
```

```
Out[94]:   S.No  Name  Age  City  Salary  Salary2
         0    1    Tom   28  Toronto   20000   141.421356
         1    2    Lee   32  HongKong    3000    54.772256
         3    4    Ram   38  Hyderabad    3900    62.449980
         2    3  Steven  43  Bay Area    8300    91.104336
```

```
In [96]: df.sort_values(by='Salary', ascending=False)
```

```
Out[96]:   S.No  Name  Age  City  Salary  Salary2
         0    1    Tom   28  Toronto   20000   141.421356
         2    3  Steven  43  Bay Area    8300    91.104336
         3    4    Ram   38  Hyderabad    3900    62.449980
         1    2    Lee   32  HongKong    3000    54.772256
```

3.5 Column Addition

```
In [21]: df=pd.read_csv("data.csv")
         df
```

```
Out[21]:   S.No  Name  Age  City  Salary
         0    1    Tom   28  Toronto   20000
         1    2    Lee   32  HongKong    3000
         2    3  Steven  43  Bay Area    8300
         3    4    Ram   38  Hyderabad    3900
```

```
In [22]: # Add new column "Address"
         df["Address"] = ["HCM", "HN", "DN", "HCM"]
         df
```

```
Out[22]:   S.No  Name  Age  City  Salary Address
         0    1    Tom   28  Toronto   20000    HCM
         1    2    Lee   32  HongKong    3000     HN
         2    3  Steven  43  Bay Area    8300     DN
         3    4    Ram   38  Hyderabad    3900    HCM
```

```
In [23]: # Or new column with default value
         df["LastSalary"] = 1000
         df
```

```
Out[23]:
```

	S.No	Name	Age	City	Salary	Address	LastSalary
0	1	Tom	28	Toronto	20000	HCM	1000
1	2	Lee	32	HongKong	3000	HN	1000
2	3	Steven	43	Bay Area	8300	DN	1000
3	4	Ram	38	Hyderabad	3900	HCM	1000

```
In [24]: # New column base on old columns
df["NewSalary"] = df.Salary + df.LastSalary
df
```

```
Out[24]:
```

	S.No	Name	Age	City	Salary	Address	LastSalary	NewSalary
0	1	Tom	28	Toronto	20000	HCM	1000	21000
1	2	Lee	32	HongKong	3000	HN	1000	4000
2	3	Steven	43	Bay Area	8300	DN	1000	9300
3	4	Ram	38	Hyderabad	3900	HCM	1000	4900

3.6 Column Deletion

Columns can be deleted or popped; let us take an example to understand how.

```
In [25]: del df["NewSalary"]
df
```

```
Out[25]:
```

	S.No	Name	Age	City	Salary	Address	LastSalary
0	1	Tom	28	Toronto	20000	HCM	1000
1	2	Lee	32	HongKong	3000	HN	1000
2	3	Steven	43	Bay Area	8300	DN	1000
3	4	Ram	38	Hyderabad	3900	HCM	1000

```
In [26]: LastSalary = df.pop('LastSalary')
df
```

```
Out[26]:
```

	S.No	Name	Age	City	Salary	Address
0	1	Tom	28	Toronto	20000	HCM
1	2	Lee	32	HongKong	3000	HN
2	3	Steven	43	Bay Area	8300	DN
3	4	Ram	38	Hyderabad	3900	HCM

```
In [27]: LastSalary
```

```
Out[27]:
```

0	1000
1	1000
2	1000
3	1000

Name: LastSalary, dtype: int64

3.7 Row Selection, Addition, and Deletion

We will now understand row selection, addition and deletion through examples

3.7.1 Selection by Label

```
In [28]: df["City"]
```

```
Out[28]: 0    Toronto
         1    HongKong
         2    Bay Area
         3    Hyderabad
         Name: City, dtype: object
```

3.7.2 Selection by integer location

```
In [29]: df.iloc[2]
```

```
Out[29]: S.No      3
         Name    Steven
         Age     43
         City    Bay Area
         Salary  8300
         Address DN
         Name: 2, dtype: object
```

3.7.3 Slice Rows

Multiple rows can be selected using ':' operator.

```
In [31]: df[2:4]
```

```
Out[31]:   S.No  Name  Age  City  Salary  Address
         2    3  Steven  43  Bay Area   8300      DN
         3    4    Ram   38  Hyderabad  3900      HCM
```

3.7.4 Addition of Rows

Add new rows to a DataFrame using the **append** function. This function will append the rows at the end.

```
In [32]: df1 = pd.DataFrame([[1, 2], [3, 4]], columns = ['a', 'b'])
         df2 = pd.DataFrame([[5, 6], [7, 8]], columns = ['a', 'b'])

         df1.append(df2)
```

```
Out[32]:   a  b
         0  1  2
         1  3  4
         0  5  6
         1  7  8
```

3.7.5 Boolean Indexing

Using a single column's values to select data.

```
In [98]: df[df.Salary > 4000]
```

```
Out[98]:
```

	S.No	Name	Age	City	Salary	Salary2
0	1	Tom	28	Toronto	20000	141.421356
2	3	Steven	43	Bay Area	8300	91.104336

Using the `isin()` method for filtering:

```
In [100]: df[df.Name.isin(['Tom', 'Ram'])]
```

```
Out[100]:
```

	S.No	Name	Age	City	Salary	Salary2
0	1	Tom	28	Toronto	20000	141.421356
3	4	Ram	38	Hyderabad	3900	62.449980

3.7.6 Drop duplicates

```
In [101]: df.drop_duplicates()
```

```
Out[101]:
```

	S.No	Name	Age	City	Salary	Salary2
0	1	Tom	28	Toronto	20000	141.421356
1	2	Lee	32	HongKong	3000	54.772256
2	3	Steven	43	Bay Area	8300	91.104336
3	4	Ram	38	Hyderabad	3900	62.449980

```
In [102]: df2 = df.copy()
df2["Company"] = [ "JVN", "JVN", "UIT", "UIT" ]
df2
```

```
Out[102]:
```

	S.No	Name	Age	City	Salary	Salary2	Company
0	1	Tom	28	Toronto	20000	141.421356	JVN
1	2	Lee	32	HongKong	3000	54.772256	JVN
2	3	Steven	43	Bay Area	8300	91.104336	UIT
3	4	Ram	38	Hyderabad	3900	62.449980	UIT

```
In [103]: df2.drop_duplicates("Company")
```

```
Out[103]:
```

	S.No	Name	Age	City	Salary	Salary2	Company
0	1	Tom	28	Toronto	20000	141.421356	JVN
2	3	Steven	43	Bay Area	8300	91.104336	UIT

4 Python Pandas - Series

Series is a one-dimensional labeled array capable of holding data of any type (integer, string, float, python objects, etc.). The axis labels are collectively called index.

4.1 Create a Series

4.1.1 Create an Empty Series

```
In [34]: s = pd.Series()  
s
```

```
Out[34]: Series([], dtype: float64)
```

4.1.2 Create a Series from ndarray

```
In [39]: # Ex1  
s = pd.Series([1,2,3,4])  
s
```

```
Out[39]: 0    1  
         1    2  
         2    3  
         3    4  
dtype: int64
```

```
In [40]: # Ex2  
import numpy as np  
data = np.array([1,2,3,4])  
s = pd.Series(data)  
s
```

```
Out[40]: 0    1  
         1    2  
         2    3  
         3    4  
dtype: int64
```

```
In [42]: # Ex3  
data = np.array(['a', 'b', 'c', 'd'])  
s = pd.Series(data, index=[100,101,102,103])  
s
```

```
Out[42]: 100    a  
         101    b  
         102    c  
         103    d  
dtype: object
```

4.1.3 Create a Series from Scalar

If data is a scalar value, an index must be provided. The value will be repeated to match the length of index

```
In [43]: s = pd.Series(5, index=[0, 1, 2, 3])  
s
```

```
Out[43]: 0    5
         1    5
         2    5
         3    5
         dtype: int64
```

4.2 Accessing Data from Series with Position

Data in the series can be accessed similar to that in an ndarray.

Ex1: Retrieve the first element. As we already know, the counting starts from zero for the array, which means the first element is stored at zeroth position and so on.

```
In [47]: s = pd.Series([1,2,3,4,5],index = ['a','b','c','d','e'])
```

```
print s

print "\nretrieve the first element:"
print s[0]
```

```
a    1
b    2
c    3
d    4
e    5
dtype: int64
```

```
retrieve the first element:
1
```

```
In [48]: # Retrieve the first three elements in the Series
print s[:3]
```

```
a    1
b    2
c    3
dtype: int64
```

4.3 Retrieve Data Using Label (Index)

A Series is like a fixed-size dict in that you can get and set values by index label.

```
In [51]: s = pd.Series([1,2,3,4,5],index = ['a','b','c','d','e'])
```

```
print s

print
```

```
#retrieve a single element
print s['a']
```

```
a    1
b    2
c    3
d    4
e    5
dtype: int64

1
```

```
In [52]: #retrieve multiple elements
print s[['a','c','d']]
```

```
a    1
c    3
d    4
dtype: int64
```

```
In [53]: # If a label is not contained, an exception is raised.
print s['f']
```

```
-----

KeyError                                Traceback (most recent call last)

<ipython-input-53-fec766af88eb> in <module>()
      1 # If a label is not contained, an exception is raised.
----> 2 print s['f']
```

```
/home/duyetdev/.local/lib/python2.7/site-packages/pandas/core/series.pyc in __getitem__(
599         key = com._apply_if_callable(key, self)
600         try:
--> 601             result = self.index.get_value(self, key)
602
603             if not is_scalar(result):

/home/duyetdev/.local/lib/python2.7/site-packages/pandas/core/indexes/base.pyc in get_value
2489             raise InvalidIndexError(key)
2490         else:
-> 2491             raise e1
2492         except Exception: # pragma: no cover
2493             raise e1
```

```
KeyError: 'f'
```

5 Python Pandas - Descriptive Statistics

A large number of methods collectively compute descriptive statistics and other related operations on DataFrame. Most of these are aggregations like `sum()`, `mean()`, but some of them, like `sumsum()`, produce an object of the same size. Generally speaking, these methods take an axis argument, just like `ndarray.{sum, std, ...}`, but the axis can be specified by name or integer.

Let's create a DataFrame and use this object throughout this chapter for all the operations.

```
In [54]: df = pd.read_csv("data.csv")
df
```

```
Out [54]:
```

	S.No	Name	Age	City	Salary
0	1	Tom	28	Toronto	20000
1	2	Lee	32	HongKong	3000
2	3	Steven	43	Bay Area	8300
3	4	Ram	38	Hyderabad	3900

5.1 `sum()`

Returns the sum of the values for the requested axis. By default, axis is index (axis=0).

```
In [55]: df.sum()
```

```
Out [55]:
```

S.No	10
Name	TomLeeStevenRam
Age	141
City	TorontoHongKongBay AreaHyderabad
Salary	35200
dtype:	object

```
In [56]: df.sum(axis=1)
```

```
Out [56]:
```

0	20029
1	3034
2	8346
3	3942
dtype:	int64

5.2 `mean()`

```
In [57]: df.mean()
```

```
Out [57]: S.No          2.50
          Age           35.25
          Salary       8800.00
          dtype: float64
```

```
In [58]: df.mean(1)
```

```
Out [58]: 0    6676.333333
          1    1011.333333
          2    2782.000000
          3    1314.000000
          dtype: float64
```

5.3 std()

Returns the Bressel standard deviation of the numerical columns.

```
In [59]: df.std()
```

```
Out [59]: S.No          1.290994
          Age           6.601767
          Salary       7817.501732
          dtype: float64
```

5.4 Functions & Description

The following table list down the important functions

S.No.	Function	Description
1	count()	Number of non-null observations
2	sum()	Sum of values
3	mean()	Mean of Values
4	median()	Median of Values
5	mode()	Mode of values
6	std()	Standard Deviation of the Values
7	min()	Minimum Value
8	max()	Maximum Value
9	abs()	Absolute Value
10	prod()	Product of Values
11	cumsum()	Cumulative Sum
12	cumprod()	Cumulative Product

5.5 Summarizing Data

The **describe()** function computes a summary of statistics pertaining to the DataFrame columns.

```
In [60]: df.describe()
```

```
Out [60]:
```

	S.No	Age	Salary
count	4.000000	4.000000	4.000000
mean	2.500000	35.250000	8800.000000
std	1.290994	6.601767	7817.501732
min	1.000000	28.000000	3000.000000
25%	1.750000	31.000000	3675.000000
50%	2.500000	35.000000	6100.000000
75%	3.250000	39.250000	11225.000000
max	4.000000	43.000000	20000.000000

And, function excludes the character columns and given summary about numeric columns. **'include'** is the argument which is used to pass necessary information regarding what columns need to be considered for summarizing. Takes the list of values; by default, 'number'.

- **object** – Summarizes String columns
- **number** – Summarizes Numeric columns
- **all** – Summarizes all columns together (Should not pass it as a list value)

```
In [62]: df.describe(include=['object'])
```

```
Out [62]:
```

	Name	City
count	4	4
unique	4	4
top	Steven	Bay Area
freq	1	1

6 Python Pandas - Function Application

To apply your own or another library's functions to Pandas objects, you should be aware of the three important methods. The methods have been discussed below. The appropriate method to use depends on whether your function expects to operate on an entire DataFrame, row- or column-wise, or element wise.

- Table wise Function Application: pipe()
- Row or Column Wise Function Application: apply()
- Element wise Function Application: applymap()

```
In [63]: df
```

```
Out [63]:
```

	S.No	Name	Age	City	Salary
0	1	Tom	28	Toronto	20000
1	2	Lee	32	HongKong	3000
2	3	Steven	43	Bay Area	8300
3	4	Ram	38	Hyderabad	3900

```
In [67]: df['Salary2'] = df.Salary.apply(np.sqrt)
df
```

```
Out[67]:
```

	S.No	Name	Age	City	Salary	Salary2
0	1	Tom	28	Toronto	20000	141.421356
1	2	Lee	32	HongKong	3000	54.772256
2	3	Steven	43	Bay Area	8300	91.104336
3	4	Ram	38	Hyderabad	3900	62.449980

7 GroupBy

Reference: https://www.tutorialspoint.com/python_pandas/python_pandas_groupby.htm

```
In [74]: df.groupby("Name").count()
```

```
Out[74]:
```

	S.No	Age	City	Salary	Salary2
Name					
Lee	1	1	1	1	1
Ram	1	1	1	1	1
Steven	1	1	1	1	1
Tom	1	1	1	1	1

```
In [75]: df.groupby("Name")['Salary'].count()
```

```
Out[75]: Name
Lee      1
Ram      1
Steven   1
Tom      1
Name: Salary, dtype: int64
```

```
In [76]: df.groupby("Name")['Salary'].mean()
```

```
Out[76]: Name
Lee      3000
Ram      3900
Steven   8300
Tom      20000
Name: Salary, dtype: int64
```

8 Merging/Joining

Pandas has full-featured, high performance in-memory join operations idiomatically very similar to relational databases like SQL.

```
pd.merge(left, right, how='inner', on=None, left_on=None, right_on=None,
         left_index=False, right_index=False, sort=True)
```

- **left** – A DataFrame object.
- **right** – Another DataFrame object.
- **on** – Columns (names) to join on. Must be found in both the left and right DataFrame objects.

- **left_on** – Columns from the left DataFrame to use as keys. Can either be column names or arrays with length equal to the length of the DataFrame.
- **right_on** – Columns from the right DataFrame to use as keys. Can either be column names or arrays with length equal to the length of the DataFrame.
- **left_index** – If True, use the index (row labels) from the left DataFrame as its join key(s). In case of a DataFrame with a MultiIndex (hierarchical), the number of levels must match the number of join keys from the right DataFrame.
- **right_index** – Same usage as left_index for the right DataFrame.
- **how** – One of 'left', 'right', 'outer', 'inner'. Defaults to inner. Each method has been described below.
- **sort** – Sort the result DataFrame by the join keys in lexicographical order. Defaults to True, setting to False will improve the performance substantially in many cases.* **left** – A DataFrame object.
- **right** – Another DataFrame object.
- **on** – Columns (names) to join on. Must be found in both the left and right DataFrame objects.
- **left_on** – Columns from the left DataFrame to use as keys. Can either be column names or arrays with length equal to the length of the DataFrame.
- **right_on** – Columns from the right DataFrame to use as keys. Can either be column names or arrays with length equal to the length of the DataFrame.
- **left_index** – If True, use the index (row labels) from the left DataFrame as its join key(s). In case of a DataFrame with a MultiIndex (hierarchical), the number of levels must match the number of join keys from the right DataFrame.
- **right_index** – Same usage as left_index for the right DataFrame.
- **how** – One of 'left', 'right', 'outer', 'inner'. Defaults to inner. Each method has been described below.
- **sort** – Sort the result DataFrame by the join keys in lexicographical order. Defaults to True, setting to False will improve the performance substantially in many cases.

```
In [78]: left = pd.DataFrame({
        'id': [1,2,3,4,5],
        'Name': ['Alex', 'Amy', 'Allen', 'Alice', 'Ayoung'],
        'subject_id': ['sub1','sub2','sub4','sub6','sub5']})
        right = pd.DataFrame(
        {'id': [1,2,3,4,5],
        'Name': ['Billy', 'Brian', 'Bran', 'Bryce', 'Betty'],
        'subject_id': ['sub2','sub4','sub3','sub6','sub5']})
        left
```

```
Out[78]:
```

	Name	id	subject_id
0	Alex	1	sub1
1	Amy	2	sub2
2	Allen	3	sub4
3	Alice	4	sub6
4	Ayoung	5	sub5

```
In [79]: right
```

```
Out[79]:
```

	Name	id	subject_id
0	Billy	1	sub2

1	Brian	2	sub4
2	Bran	3	sub3
3	Bryce	4	sub6
4	Betty	5	sub5

8.1 Merge Two DataFrames on a Key

```
In [80]: pd.merge(left, right, on='id')
```

```
Out[80]:
```

	Name_x	id	subject_id_x	Name_y	subject_id_y
0	Alex	1	sub1	Billy	sub2
1	Amy	2	sub2	Brian	sub4
2	Allen	3	sub4	Bran	sub3
3	Alice	4	sub6	Bryce	sub6
4	Ayoung	5	sub5	Betty	sub5

8.2 Merge Two DataFrames on Multiple Keys

```
In [81]: pd.merge(left, right, on=['id', 'subject_id'])
```

```
Out[81]:
```

	Name_x	id	subject_id	Name_y
0	Alice	4	sub6	Bryce
1	Ayoung	5	sub5	Betty

8.3 Left/Right Join

```
In [82]: pd.merge(left, right, on='subject_id', how='left')
```

```
Out[82]:
```

	Name_x	id_x	subject_id	Name_y	id_y
0	Alex	1	sub1	NaN	NaN
1	Amy	2	sub2	Billy	1.0
2	Allen	3	sub4	Brian	2.0
3	Alice	4	sub6	Bryce	4.0
4	Ayoung	5	sub5	Betty	5.0

```
In [83]: pd.merge(left, right, on='subject_id', how='right')
```

```
Out[83]:
```

	Name_x	id_x	subject_id	Name_y	id_y
0	Amy	2.0	sub2	Billy	1
1	Allen	3.0	sub4	Brian	2
2	Alice	4.0	sub6	Bryce	4
3	Ayoung	5.0	sub5	Betty	5
4	NaN	NaN	sub3	Bran	3

8.4 Inner/Outer Join

```
In [84]: pd.merge(left, right, on='subject_id', how='inner')
```

```
Out[84]:
```

	Name_x	id_x	subject_id	Name_y	id_y
0	Amy	2	sub2	Billy	1
1	Allen	3	sub4	Brian	2
2	Alice	4	sub6	Bryce	4
3	Ayoung	5	sub5	Betty	5

```
In [85]: pd.merge(left, right, on='subject_id', how='outer')
```

```
Out[85]:
```

	Name_x	id_x	subject_id	Name_y	id_y
0	Alex	1.0	sub1	NaN	NaN
1	Amy	2.0	sub2	Billy	1.0
2	Allen	3.0	sub4	Brian	2.0
3	Alice	4.0	sub6	Bryce	4.0
4	Ayoung	5.0	sub5	Betty	5.0
5	NaN	NaN	sub3	Bran	3.0

9 Concatenation

Pandas provides various facilities for easily combining together Series, DataFrame, and Panel objects.

```
pd.concat(objs,axis=0,join='outer',join_axes=None, ignore_index=False)
```

```
In [86]: pd.concat([left, right])
```

```
Out[86]:
```

	Name	id	subject_id
0	Alex	1	sub1
1	Amy	2	sub2
2	Allen	3	sub4
3	Alice	4	sub6
4	Ayoung	5	sub5
0	Billy	1	sub2
1	Brian	2	sub4
2	Bran	3	sub3
3	Bryce	4	sub6
4	Betty	5	sub5

```
In [87]: pd.concat([left, right], keys=['x', 'y'])
```

```
Out[87]:
```

		Name	id	subject_id
x	0	Alex	1	sub1
	1	Amy	2	sub2
	2	Allen	3	sub4
	3	Alice	4	sub6
	4	Ayoung	5	sub5
y	0	Billy	1	sub2
	1	Brian	2	sub4
	2	Bran	3	sub3
	3	Bryce	4	sub6
	4	Betty	5	sub5

Concatenating Using append

```
In [88]: left.append(right)
```

```
Out[88]:
```

	Name	id	subject_id
0	Alex	1	sub1
1	Amy	2	sub2
2	Allen	3	sub4
3	Alice	4	sub6
4	Ayoung	5	sub5
0	Billy	1	sub2
1	Brian	2	sub4
2	Bran	3	sub3
3	Bryce	4	sub6
4	Betty	5	sub5

10 Histogramming

See more at [Histogramming and Discretization](#)

```
In [108]: s = pd.Series(np.random.randint(0, 7, size=10))
```

```
In [109]: s.value_counts()
```

```
Out[109]: 6    3
          5    2
          4    1
          3    1
          2    1
          1    1
          0    1
          dtype: int64
```

11 Time Series

pandas has simple, powerful, and efficient functionality for performing resampling operations during frequency conversion (e.g., converting secondly data into 5-minutely data). This is extremely common in, but not limited to, financial applications. See the [Time Series section](#)

```
In [112]: rng = pd.date_range('1/1/2012', periods=100, freq='S')
          rng[:5]
```

```
Out[112]: DatetimeIndex(['2012-01-01 00:00:00', '2012-01-01 00:00:01',
                          '2012-01-01 00:00:02', '2012-01-01 00:00:03',
                          '2012-01-01 00:00:04'],
                          dtype='datetime64[ns]', freq='S')
```

```
In [113]: ts = pd.Series(np.random.randint(0, 500, len(rng)), index=rng)
          ts
```

```

Out[113]: 2012-01-01 00:00:00    187
          2012-01-01 00:00:01    314
          2012-01-01 00:00:02    211
          2012-01-01 00:00:03    141
          2012-01-01 00:00:04     90
          2012-01-01 00:00:05      0
          2012-01-01 00:00:06    255
          2012-01-01 00:00:07    393
          2012-01-01 00:00:08    237
          2012-01-01 00:00:09    109
          2012-01-01 00:00:10     90
          2012-01-01 00:00:11    115
          2012-01-01 00:00:12     91
          2012-01-01 00:00:13      7
          2012-01-01 00:00:14    472
          2012-01-01 00:00:15    125
          2012-01-01 00:00:16    468
          2012-01-01 00:00:17    243
          2012-01-01 00:00:18    495
          2012-01-01 00:00:19    221
          2012-01-01 00:00:20    428
          2012-01-01 00:00:21     35
          2012-01-01 00:00:22    451
          2012-01-01 00:00:23    440
          2012-01-01 00:00:24    213
          2012-01-01 00:00:25    171
          2012-01-01 00:00:26    267
          2012-01-01 00:00:27    167
          2012-01-01 00:00:28    484
          2012-01-01 00:00:29    454
          ...
          2012-01-01 00:01:10     33
          2012-01-01 00:01:11    209
          2012-01-01 00:01:12     37
          2012-01-01 00:01:13    495
          2012-01-01 00:01:14     46
          2012-01-01 00:01:15    440
          2012-01-01 00:01:16    349
          2012-01-01 00:01:17    170
          2012-01-01 00:01:18    353
          2012-01-01 00:01:19    172
          2012-01-01 00:01:20     76
          2012-01-01 00:01:21    475
          2012-01-01 00:01:22    360
          2012-01-01 00:01:23    185
          2012-01-01 00:01:24    305
          2012-01-01 00:01:25    458
          2012-01-01 00:01:26    222

```

```
2012-01-01 00:01:27    264
2012-01-01 00:01:28    480
2012-01-01 00:01:29    418
2012-01-01 00:01:30    255
2012-01-01 00:01:31    293
2012-01-01 00:01:32    146
2012-01-01 00:01:33    199
2012-01-01 00:01:34     92
2012-01-01 00:01:35    305
2012-01-01 00:01:36    454
2012-01-01 00:01:37     45
2012-01-01 00:01:38     71
2012-01-01 00:01:39    436
Freq: S, Length: 100, dtype: int64
```

```
In [114]: ts.resample('5Min').sum()
```


```
Out[114]: 2012-01-01    26667
Freq: 5T, dtype: int64
```

12 Plotting

See more: <https://pandas.pydata.org/pandas-docs/stable/visualization.html#visualization>


```
In [5]: ts = pd.Series(np.random.randn(1000), index=pd.date_range('1/1/2000', periods=1000))
        ts = ts.cumsum()
        ts.plot()
```

```
Out[5]: <matplotlib.axes._subplots.AxesSubplot at 0x7f73e36a7710>
```




```
In [8]: df.Age.plot(kind='bar')
```

```
Out[8]: <matplotlib.axes._subplots.AxesSubplot at 0x7f73e12e6390>
```




```
In [29]: df.Salary.plot.bar(alpha=0.5)
```

```
Out[29]: <matplotlib.axes._subplots.AxesSubplot at 0x7f73dfb29450>
```



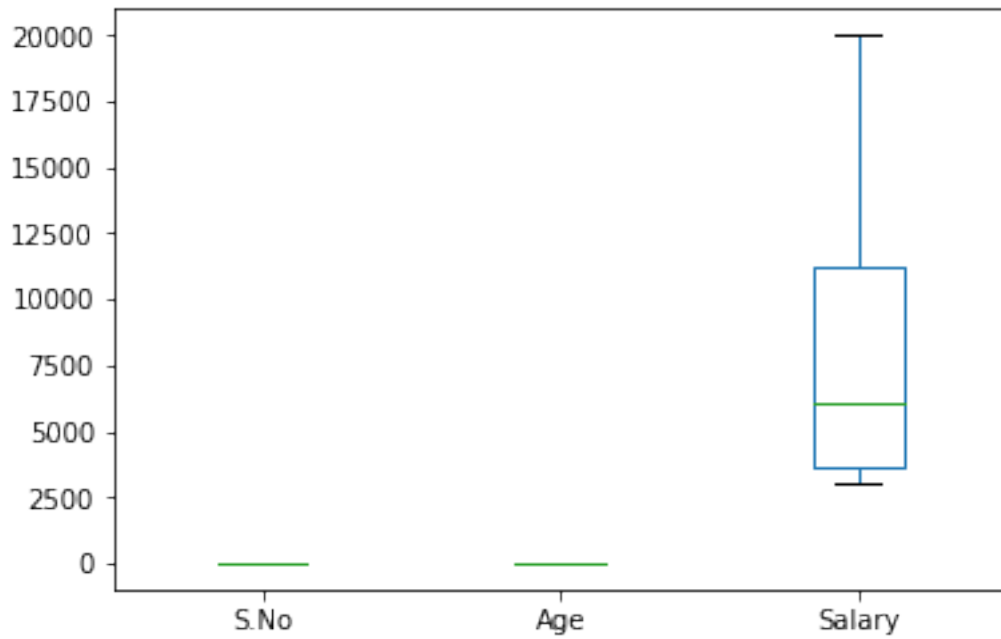
```
In [13]: df.plot.hist(alpha=0.5)
```

```
Out[13]: <matplotlib.axes._subplots.AxesSubplot at 0x7f73e0f27cd0>
```



```
In [14]: df.plot.box()
```

```
Out[14]: <matplotlib.axes._subplots.AxesSubplot at 0x7f73e0c74190>
```




```
In [17]: df.boxplot(by='S.No')
```

```
Out[17]: array([<matplotlib.axes._subplots.AxesSubplot object at 0x7f73e066d3d0>,  
                <matplotlib.axes._subplots.AxesSubplot object at 0x7f73e05e4c50>], dtype=object)
```




```
In [30]: from pandas.plotting import scatter_matrix
df_plot = pd.DataFrame(np.random.randn(1000, 4), columns=['a', 'b', 'c', 'd'])
scatter_matrix(df_plot, alpha=0.2, figsize=(6, 6), diagonal='kde')
```

```
Out[30]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x7f73d70c9f10>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7f73d605e810>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7f73d5f65890>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7f73d5f57310>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x7f73d5ede350>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7f73d5e43c90>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7f73d5dcae90>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7f73d5d40750>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x7f73d5cc6850>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7f73d5c392d0>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7f73d5bbe750>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7f73d5bf3f90>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x7f73d5b5aa10>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7f73d5ae0890>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7f73d5a46f50>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7f73d59cdf90>]], dtype=object)
```



13 References

- [Python Pandas Tutorial \(tutorialspoint\)](#)
- [Lessons for New pandas Users \(pandas.pydata.org\)](#)
- [Exercises for New Users \(pandas.pydata.org\)](#)